

Shantanu Joshi

J073

In [2]:

```
import os
print(os.getcwd())
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

C:\Users\shantanu\Desktop\Important

In [3]:

```
df = pd.read_csv('C:/Users/shantanu/Downloads/car_evaluation.csv', header = None)
```

In [4]:

```
df.head()
```

Out[4]:

	0	1	2	3	4	5	6
0	vhigh	vhigh	2	2	small	low	unacc
1	vhigh	vhigh	2	2	small	med	unacc
2	vhigh	vhigh	2	2	small	high	unacc
3	vhigh	vhigh	2	2	med	low	unacc
4	vhigh	vhigh	2	2	med	med	unacc

In [5]:

```
col_names = ['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety', 'class']
df.columns = col_names
col_names
```

Out[5]:

```
['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety', 'class']
```

In [6]:

```
df.head()
```

Out[6]:

	buying	maint	doors	persons	lug_boot	safety	class
0	vhigh	vhigh	2	2	small	low	unacc
1	vhigh	vhigh	2	2	small	med	unacc
2	vhigh	vhigh	2	2	small	high	unacc
3	vhigh	vhigh	2	2	med	low	unacc
4	vhigh	vhigh	2	2	med	med	unacc

In [7]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1728 entries, 0 to 1727
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   buying      1728 non-null   object
 1   maint       1728 non-null   object
 2   doors       1728 non-null   object
 3   persons     1728 non-null   object
 4   lug_boot    1728 non-null   object
 5   safety      1728 non-null   object
 6   class       1728 non-null   object
dtypes: object(7)
memory usage: 94.6+ KB
```

In [8]:

```
for i in col_names:  
    print(df[i].value_counts())
```

```
vhigh    432  
high     432  
med      432  
low      432  
Name: buying, dtype: int64  
vhigh    432  
high     432  
med      432  
low      432  
Name: maint, dtype: int64  
2         432  
3         432  
4         432  
5more     432  
Name: doors, dtype: int64  
2         576  
4         576  
more      576  
Name: persons, dtype: int64  
small     576  
med       576  
big       576  
Name: lug_boot, dtype: int64  
low       576  
med       576  
high      576  
Name: safety, dtype: int64  
unacc    1210  
acc       384  
good       69  
vgood      65  
Name: class, dtype: int64
```

In [9]:

```
df.shape
```

Out[9]:

```
(1728, 7)
```

In [10]:

```
X = df.drop(['class'],axis = 1)  
y = df['class']
```

In [11]:

```
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.3,random_state=42)
```

In [12]:

```
from sklearn.preprocessing import OrdinalEncoder
enc = OrdinalEncoder()
X_train = enc.fit_transform(X_train)
X_test = enc.transform(X_test)
```

Gini index as criterion

In [13]:

```
from sklearn.tree import DecisionTreeClassifier
```

In [14]:

```
clf_gini = DecisionTreeClassifier(criterion='gini', max_depth=3, random_state=42)
clf_gini.fit(X_train, y_train)
```

Out[14]:

```
DecisionTreeClassifier(max_depth=3, random_state=42)
```

In [15]:

```
y_pred = clf_gini.predict(X_test)
```

Grid Search Cv

In [16]:

```
from sklearn.model_selection import GridSearchCV
option=['gini','entropy']
weight_option=['auto','sqrt','log2']
param_grid = {'criterion': option , 'max_features':[2,3,4,5,6] , 'max_depth':[4,5,6,7] , 'm
grid=GridSearchCV(clf_gini,param_grid,cv=3,scoring='accuracy')
grid.fit(X_train,y_train)
print(grid.best_score_)
print(grid.best_params_)
```

```
0.9247311827956989
```

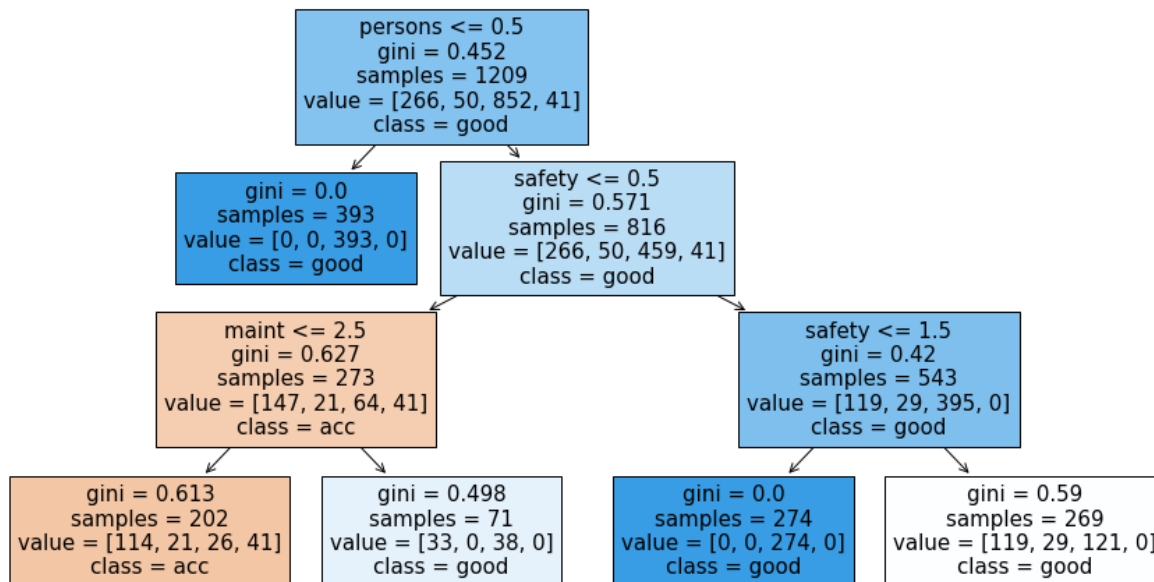
```
{'criterion': 'gini', 'max_depth': 7, 'max_features': 6, 'min_samples_split': 2}
```

In [17]:

```

from sklearn import tree
plt.figure(figsize=(15,8))
tree.plot_tree(clf_gini,
               feature_names=['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety'],
               class_names= list(set(y_train)),
               filled = True)
plt.show()

```



In [18]:

```

# Check for underfitting

print(f'Training set score: {clf_gini.score(X_train,y_train)}')
print(f'Test set score: {clf_gini.score(X_test,y_test)}')

```

Training set score: 0.7775020678246485

Test set score: 0.7572254335260116

Model after grid search

In [19]:

```
dtc = DecisionTreeClassifier(criterion='gini', max_depth=7,max_features = 6)
dtc.fit(X_train, y_train)
```

Out[19]:

```
DecisionTreeClassifier(max_depth=7, max_features=6)
```

In [20]:

```
y_pred = dtc.predict(X_test)
```

In [21]:

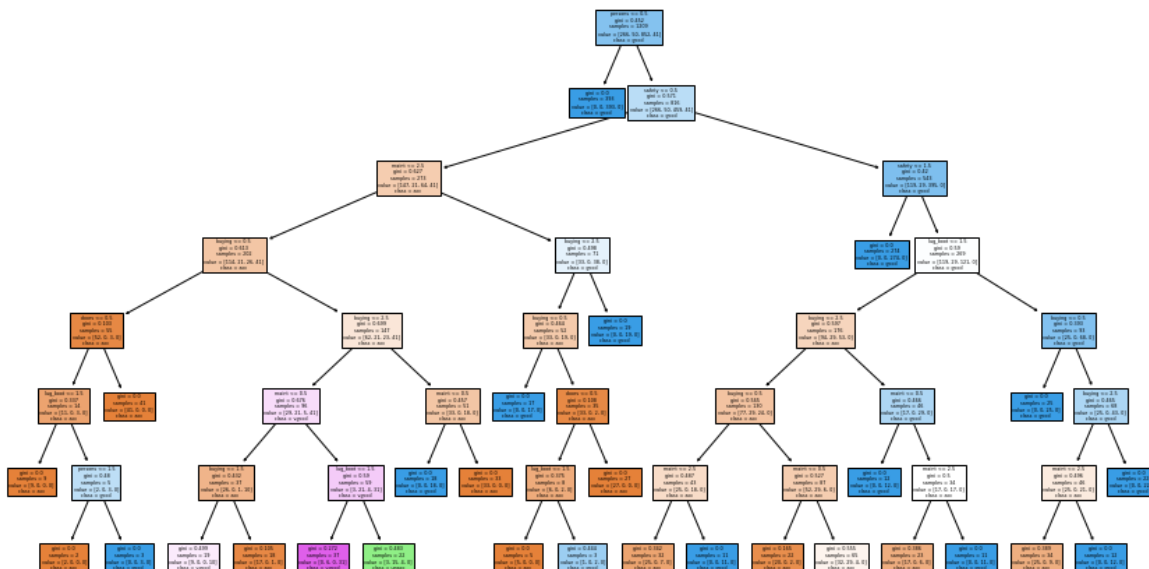
```
print(f'Training set score: {dtc.score(X_train,y_train)}')
print(f'Test set score: {dtc.score(X_test,y_test)}')
```

```
Training set score: 0.9330024813895782
```

```
Test set score: 0.9344894026974951
```

In [22]:

```
from sklearn import tree
plt.figure(figsize=(15,8))
tree.plot_tree(dtc,
               feature_names=['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety'],
               class_names= list(set(y_train)),
               filled = True)
plt.show()
```



Cross validation

In [23]:

```
from sklearn.model_selection import cross_val_score
score=cross_val_score(dtc,X_train,y_train,cv=10,scoring='accuracy')
score.mean()
```

Out[23]:

0.920564738292011

In [24]:

```
from sklearn.model_selection import cross_val_score
score=cross_val_score(dtc,X_test,y_test,cv=10,scoring='accuracy')
score.mean()
```

Out[24]:

0.8978883861236803

In [25]:

```
from sklearn.metrics import confusion_matrix, classification_report
cm = confusion_matrix(y_test, y_pred)
```

In [26]:

```
print(cm)
```

```
[[109  4  1  4]
 [ 10  6  0  3]
 [ 11  0 346  1]
 [  0  0  0 24]]
```

In [28]:

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
acc	0.84	0.92	0.88	118
good	0.60	0.32	0.41	19
unacc	1.00	0.97	0.98	358
vgood	0.75	1.00	0.86	24
accuracy			0.93	519
macro avg	0.80	0.80	0.78	519
weighted avg	0.94	0.93	0.93	519

In []:

